Observations

Lab note for Retry Loom for tool failures.

Multi-step runs were aborting when a single retrieval call timed out.

The observations follow.

Retry Loom repeats a failed tool call once with a narrower query before it returns a partial brief.

Capped retries cut abandoned runs from 11 percent to 3 percent on the operations queue.

The working decision is: Adopt a single retry with a narrower query, then finish with whatever evidence is already in hand..

Measurement

How it was measured: The queue replay compared a no-retry agent with a one-retry agent on 300 sessions that included injected timeouts.

What came back: Completed sessions rose from 89 percent to 97 percent and median added delay was 1.4 seconds.

What this does not show: The retry policy was only exercised on retrieval timeouts, not on malformed model outputs.

Still unresolved: Should a second failure switch models or simply stop with the evidence collected so far? Operational risks: Severity high: Unbounded retries can loop on a tool that fails for a permanent reason such as a missing file.

Recommended action: Allow at most one retry and record the tool error in the trajectory..