Observations

Lab note for LoRA Ridge rank sweep.

A rank sweep asked how much adapter capacity the brief task actually needed.

The observations follow.

Raising LoRA rank from 8 to 16 improved structured field F1, while rank 32 did not improve it further.

Rank 32 increased trainable parameters by about 2 times versus rank 16 with no gain on the held-out briefs.

The working decision is: Use rank 16 for the brief adapter unless a later mixture shows a clear gain from more capacity..

Measurement

How it was measured: Ranks 8, 16, and 32 were trained for one epoch on the same examples and scored on documents withheld by identifier.

What came back: Structured field F1 was 0.62 at rank 8, 0.74 at rank 16, and 0.73 at rank 32.

What this does not show: The sweep used one base model and one task mix, so the plateau may sit elsewhere for other bases.

Still unresolved: Does a larger rank help only when the training mixture includes comparison as well as extraction? Operational risks: Severity medium: Training a higher rank on a small corpus can memorize citation strings from the training notes.

Recommended action: Keep entire documents, not individual examples, in the held-out split..