Summary

Memorandum on Judge Caliper preference checks.

Automatic preference labels were drifting away from what analysts actually chose.

Judge Caliper treats a tuned answer as a win only when its field overlap exceeds the base answer by more than two points.

Ties inside that band prevented a noisy judge from inflating the win rate.

Operator decision: Publish both win rates and keep a two-point margin before calling a preference..

Evidence

Method note: Analyst choices on 40 pairs were compared with caliper labels that used a two-point field-overlap margin.

Reported result: Agreement with analysts was 0.85 when ties were set aside and 0.78 when ties counted as non-wins.

Limits of the evidence: The caliper was computed on structured fields and does not capture prose clarity.

Open questions: Would analysts still agree with the caliper on briefs that are equally complete but differently ordered? Risk review: Severity low: Counting ties as wins would overstate the value of a model that mostly matches the baseline.

Recommended action: Report win rate both with ties excluded and with ties included in the denominator..